

Your AI Is Lying to You

13 techniques to make it stop.

by d.avebrown

These models are word predictors — fancy ones, but at their core, they're guessing the next token based on patterns in their training data. And they've been trained to be helpful at all costs. To agree. To sound confident.

Which is great until they don't know the answer. Then they make one up. Smoothly. Convincingly. Completely wrong.

Here are 13 field-tested techniques to catch it, prevent it, and get dramatically better outputs from any AI model.

FROM THE VIDEO

The three fundamentals

01 — Tell It Not to Lie

Sounds almost too simple. But these models are trained to follow instructions — and most people never think to give this one. A single line in your prompt changes the entire dynamic.

TRY THIS

"Only give me answers you can back up with real sources. If you're not sure, say so."

02 — Tell It to Search, Not Guess

If your question involves niche topics or anything that happened recently, the model probably doesn't know. Don't let it wing it. Force it to look things up instead of generating from memory.

TRY THIS

"Search for the latest information on [topic]. Do not answer from memory."

03 — Break the Echo Chamber

Long chat sessions make your AI dumber. The deeper you go in one conversation, the more it agrees with itself — and with you. Copy your output, start fresh, and attack it from a completely different angle.

TRY THIS

"A junior engineer wrote this. Tear it apart."

BONUS TECHNIQUES

Ten more ways to keep your AI honest

04 — Tell It the Stakes

Models calibrate caution based on context. Higher stakes = more hedging, more "verify this" flags, fewer confident guesses.

TRY THIS

"This information will be used to make a real business decision." or "A doctor will act on this."

05 — Run the Confidence Check

Ask it to rate its own certainty on each claim. Watch how fast it starts hedging. The claims it's least confident about are exactly the ones you need to verify independently.

TRY THIS

"On a scale of 1–10, how confident are you in each of these claims?"

06 — Give It Permission to Say "I Don't Know"

Models hallucinate partly because they think silence equals failure. They'd rather be confidently wrong than admit uncertainty. Give them an explicit out and the output quality jumps.

TRY THIS

"If you don't know, just say 'I don't know' — that's the correct answer here."

07 — Separate Tasks from Questions

Don't combine research and writing in one prompt. "Write a bio and include her career history" asks the model to fact-find and generate simultaneously. Hallucinations love hiding inside creation tasks. Get the facts first, confirm them, then write.

08 — Interrogate the Output

Instead of fact-checking every line yourself, make the model do it. Ask it to poke holes in its own reasoning. It's surprisingly good at this — you just have to trigger it.

TRY THIS

"What would have to be true for this to be wrong?"

09 — Use Two Models, Not One

Run the same question through Claude and ChatGPT. Where they agree, you're probably fine. Where they disagree is your uncertainty zone — don't trust either answer there without checking a primary source.

10 — Anchor It to Primary Sources

Paste in the actual document, data, or article. This turns the model from a guesser into a reader — and readers don't hallucinate. The more real context you feed it, the less it needs to invent.

TRY THIS

"Answer only from what I've given you. Do not bring in outside information."

11 — Name the Expert

"Answer as an expert" activates broad, generic knowledge. A specific role activates careful, precise reasoning. The more detailed the persona you assign, the less room there is for slop.

TRY THIS

"Answer as a forensic accountant who's testified in federal court."

12 — Break the Sycophancy Loop

Without an explicit trigger to push back, the model will validate whatever framing you brought in. It's designed to agree with you. You have to deliberately force it to argue the other side.

TRY THIS

"Assume I'm wrong about my core assumption. What's the argument against it?"

13 — Version Your Prompts

If you're running the same type of query regularly, write your anti-hallucination instructions once and reuse them every time. Don't rely on the model being in a good mood. Build a prompt template with your safeguards baked in. Consistency beats hope.

Want more AI tips that actually work?

Follow me on Instagram for weekly breakdowns.

@d.avebrown